

ass2_grad

April 1, 2018

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from data_util_KNN import readMnist

In [2]: X_train, y_train, X_test, y_test = readMnist('E:\Code\jupyterpy\Digital'
                                                    'Video\datasets\mnist')

In [3]: mask = range(5000)
X_tr = X_train[mask]
y_tr = y_train[mask]
mask = range(500)
X_te = X_test[mask]
y_te = y_test[mask]

In [4]: W = np.random.random((X_tr.shape[1] + 1, 10))
print(W.shape)
W

(785, 10)

Out[4]: array([[ 0.80494852,  0.90765804,  0.72454771, ...,  0.30158744,
                 0.90272307,  0.6233478 ],
               [ 0.91114889,  0.84165155,  0.70317371, ...,  0.5170203 ,
                 0.56772921,  0.09373554],
               [ 0.66437615,  0.66651792,  0.61080153, ...,  0.56486989,
                 0.79127343,  0.17587681],
               ...,
               [ 0.45672482,  0.14032603,  0.73310886, ...,  0.57979093,
                 0.28008893,  0.80227858],
               [ 0.76021803,  0.60730959,  0.32624087, ...,  0.57817323,
                 0.29905183,  0.71498281],
               [ 0.62510794,  0.48628265,  0.45963862, ...,  0.56807162,
                 0.22437052,  0.32629691]])

In [5]: def featurenormal(X):
    std = np.std(X, axis=1).reshape(-1, 1)
    mean = np.mean(X, axis=1).reshape(-1, 1)
```

```

X_change = (X - mean) / std
return X_change, mean, std
X_tr, mean, std= featurenormal(X_tr)

```

#softmaxLoss

$$L_i = -\log \frac{e^{s_k}}{\sum_j e^{s_j}}$$

$$L_i = -\sum_k P_{ik} \log P_k \quad \frac{\partial L_i}{\partial P_k} = -\frac{1}{P_k}$$

$$P_k = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

$$k = m : \quad \frac{\partial P_k}{\partial s_m} = -\frac{e^{s_m} \sum_j e^{s_j} - (e^{s_m})^2}{(\sum_j e^{s_j})^2} = P_m(1 - P_m) \quad k \neq m : \quad \frac{\partial P_k}{\partial s_m} = -\frac{e^{s_k} e^{s_m}}{(\sum_j e^{s_j})^2} = -P_k P_m$$

$$s_m = (X_i \times W)_m \quad \frac{\partial s_m}{\partial W} = X_i$$

$$k = m : \quad \frac{\partial L_i}{\partial W} = P_m - 1 \quad k \neq m : \quad \frac{\partial L_i}{\partial W} = P_m$$

```

In [6]: def computeSoftmaxLoss(X_tr, y_tr, W, reg):
    n, w = X_tr.shape
    X_tr = np.c_[np.ones((n, 1)), X_tr]
    s_mat = X_tr.dot(W)
    s_m_argmax = np.argmax(s_mat, axis=1)
    s_m_max = s_mat[range(n), s_m_argmax].reshape(-1, 1)
    s_m_exp = np.exp(s_mat - s_m_max)
    L_mat = - np.log(s_m_exp[range(n), y_tr] / np.sum(s_m_exp, axis=1))
    L = np.sum(L_mat) / n
    L += 0.5 * reg * np.sum(W ** 2)
    dW = np.zeros(W.shape)

    '''
    for i in range(n):
        for m in range(W.shape[1]):
            P_m = np.reshape((s_m_exp[range(n), m] /
                               np.sum(s_m_exp, axis=1)), (-1, 1))
            if (m == y_tr[i]):
                dW[:, m] += X_tr[i] * (P_m[i] - 1)
            else:
                dW[:, m] += X_tr[i] * (P_m[i])
    '''

    P_m = (s_m_exp / np.sum(s_m_exp, axis=1).reshape(-1, 1))
    P_ik = np.zeros(P_m.shape)
    P_ik[range(n), y_tr] = 1

```

```

dW = X_tr.T.dot(P_m - P_ik)
dW /= n
dW += reg * W
return L, dW

```

```

In [7]: loss_softmax, grad_softmax = computeSoftmaxLoss(X_tr[:10], y_tr[:10],
                                                         W, reg=0)

```

```

In [8]: def numComputeSoftmax(X_tr, y_tr, W, reg, h=0.000001):
    w, k = W.shape
    dW = np.zeros((w, k))
    for i in range(w):
        for j in range(k):
            loss = computeSoftmaxLoss(X_tr[:10], y_tr[:10], W, 0)[0]
            W[i, j] += h
            loss_c = computeSoftmaxLoss(X_tr[:10], y_tr[:10], W, 0)[0]
            dW[i, j] = (loss_c - loss) / h
    return dW

```

```

In [9]: num_softmax = numComputeSoftmax(X_tr[:10], y_tr[:10], W, 0)
num_softmax

```

```

Out[9]: array([[ -2.85524191e-02,  -2.93491350e-01,   9.50748813e-04, ...,
                 1.45805453e-01,   3.43865381e-04,  -9.94415448e-02],
               [  1.57605609e-02,   9.00941028e-02,   1.29820314e-02, ...,
                 -6.60884787e-02,  -1.24215305e-04,   3.96309474e-02],
               [  1.57605626e-02,   9.00941011e-02,   1.29820314e-02, ...,
                 -6.60884805e-02,  -1.24215305e-04,   3.96309456e-02],
               ...,
               [  1.57605644e-02,   9.00941011e-02,   1.29820314e-02, ...,
                 -6.60884787e-02,  -1.24215305e-04,   3.96309456e-02],
               [  1.57605609e-02,   9.00941046e-02,   1.29820314e-02, ...,
                 -6.60884787e-02,  -1.24215305e-04,   3.96309456e-02],
               [  1.57605626e-02,   9.00941046e-02,   1.29820297e-02, ...,
                 -6.60884787e-02,  -1.24215305e-04,   3.96309456e-02]])

```

```

In [10]: grad_softmax

```

```

Out[10]: array([[ -2.85524313e-02,  -2.93491353e-01,   9.50748654e-04, ...,
                 1.45805473e-01,   3.43866960e-04,  -9.94415454e-02],
               [  1.57605607e-02,   9.00941018e-02,   1.29820311e-02, ...,
                 -6.60884772e-02,  -1.24215382e-04,   3.96309468e-02],
               [  1.57605607e-02,   9.00941018e-02,   1.29820311e-02, ...,
                 -6.60884772e-02,  -1.24215382e-04,   3.96309468e-02],
               ...,
               [  1.57605607e-02,   9.00941018e-02,   1.29820311e-02, ...,
                 -6.60884772e-02,  -1.24215382e-04,   3.96309468e-02],
               [  1.57605607e-02,   9.00941018e-02,   1.29820311e-02, ...,
                 -6.60884772e-02,  -1.24215382e-04,   3.96309468e-02],

```

```
[ 1.57605607e-02,  9.00941018e-02,  1.29820311e-02, ...,
 -6.60884772e-02, -1.24215382e-04,  3.96309468e-02]])
```

```
In [11]: diff = np.linalg.norm(num_softmax - grad_softmax) / \
           np.linalg.norm(num_softmax + grad_softmax)
           diff
```

```
Out[11]: 8.9204458062729595e-08
```

##HingeLoss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

$$k = m: \quad L_i = \sum_{j \neq k} \max(0, s_j - s_k + 1) \quad \frac{\partial L_i}{\partial s_m} = \sum_{j \neq k} -1 \quad \text{if} : (s_j - s_k + 1) > 0$$

$$k \neq m: \quad L_i = \sum_{j \neq k} \max(0, s_j - s_k + 1) \quad \frac{\partial L_i}{\partial s_m} = \sum_{j \neq k} 1 \quad \text{if} : (s_j - s_k + 1) > 0$$

$$s_m = (X_i \times W)_m \quad \frac{\partial s_m}{\partial W} = X_i$$

$$k = m: \quad \frac{\partial L_i}{\partial W} = \sum_{j \neq k} -X_i \quad \text{if} : (s_j - s_k + 1) > 0 \quad k \neq m: \quad \frac{\partial L_i}{\partial W} = \sum_{j \neq k} X_i \quad \text{if} : (s_j - s_k + 1) > 0$$

```
In [12]: def computeHingeLoss(X_tr, y_tr, W, reg):
           n, w = X_tr.shape
           X_tr = np.c_[np.ones((n, 1)), X_tr]
           s_mat = X_tr.dot(W)
           s_k = s_mat[range(n), y_tr]
           pred_score = s_mat - s_k.reshape(-1, 1) + 1
           pred_score[range(n), y_tr] = 0
           L = np.sum(pred_score[np.where(pred_score > 0)]) / n
           L += 0.5 * reg * np.sum(W ** 2)
           dW = np.zeros(W.shape)
           pred_score[pred_score < 0] = 0
           '''
           for i in range(n):
               for m in range(W.shape[1]):
                   if pred_score[i, m] > 0:
                       dW[:, m] += X_tr[i]
                       dW[:, y_tr[i]] -= X_tr[i]
           '''

           pred_score[pred_score > 0] = 1
           sum_score = np.sum(pred_score, axis=1)
           pred_score[range(n), y_tr] = - sum_score
           dW = X_tr.T.dot(pred_score)
```

```

dW /= n
dW += reg * W
return L, dW

```

```
In [13]: loss_hinge, grad_hinge = computeHingeLoss(X_tr[:10], y_tr[:10], W, 0)
```

```
In [14]: def numHingeCompute(X_tr, y_tr, W, reg, h=0.000001):
    w, k = W.shape
    dW = np.zeros((w, k))
    for i in range(w):
        for j in range(k):
            loss = computeHingeLoss(X_tr[:10], y_tr[:10], W, 0)[0]
            W[i, j] += h
            loss_c = computeHingeLoss(X_tr[:10], y_tr[:10], W, 0)[0]
            dW[i, j] = (loss_c - loss) / h
    return dW

```

```
In [15]: num_hinge = numHingeCompute(X_tr[:10], y_tr[:10], W, 0)
num_hinge
```

```
Out[15]: array([[ -0.1          , -1.10000001, -0.09999999, ...,  0.8          ,
                  0.3          , -0.50000001],
 [ 0.05048132,  0.31159955,  0.0716095 , ..., -0.32712889,
 -0.11203351,  0.21915552],
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712889,
 -0.1120335 ,  0.21915551],
 ...,
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712888,
 -0.11203351,  0.21915552],
 [ 0.05048132,  0.31159955,  0.0716095 , ..., -0.32712888,
 -0.11203351,  0.21915551],
 [ 0.05048132,  0.31159955,  0.0716095 , ..., -0.32712888,
 -0.11203351,  0.21915551]])

```

```
In [16]: grad_hinge
```

```
Out[16]: array([[ -0.1          , -1.1          , -0.1          , ...,  0.8          ,
                  0.3          , -0.5          ],
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712889,
 -0.1120335 ,  0.21915551],
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712889,
 -0.1120335 ,  0.21915551],
 ...,
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712889,
 -0.1120335 ,  0.21915551],
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712889,
 -0.1120335 ,  0.21915551],
 [ 0.05048132,  0.31159955,  0.07160951, ..., -0.32712889,
 -0.1120335 ,  0.21915551]])

```

```
In [17]: diff = np.linalg.norm(num_hinge - grad_hinge) / \
           np.linalg.norm(num_hinge + grad_hinge)
           diff
```

```
Out[17]: 4.012536292715479e-09
```